



Hardware deceleration

The challenges of speeding up software

Intel Programmable Solutions Group

Hardware acceleration - what is the challenge?

FPGA and SoC FPGA devices can be used as hardware accelerators to software

- Improve performance of calculations
- Improve power efficiency
- Can often be programmable to adapt to workloads at runtime

How do we ensure the right architecture to achieve better performance?
What are the pitfalls if we get the architecture wrong?

If we get the architecture wrong, then we could end up with
Hardware Deceleration!

Before trying to accelerate with hardware....

Let's ensure the software is optimized first




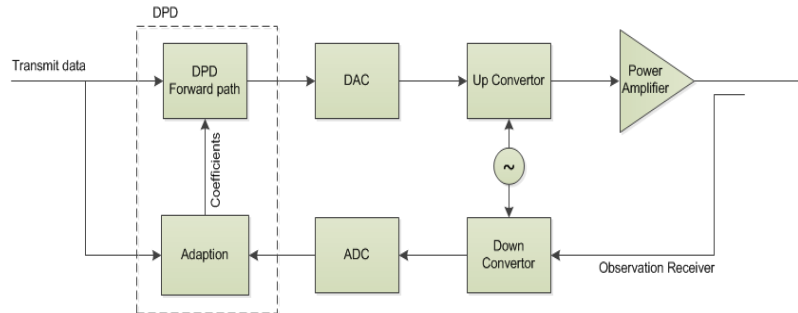
Example - DPD - Digital Predistortion

Power amplifiers (PA) for communications are not linear in their output


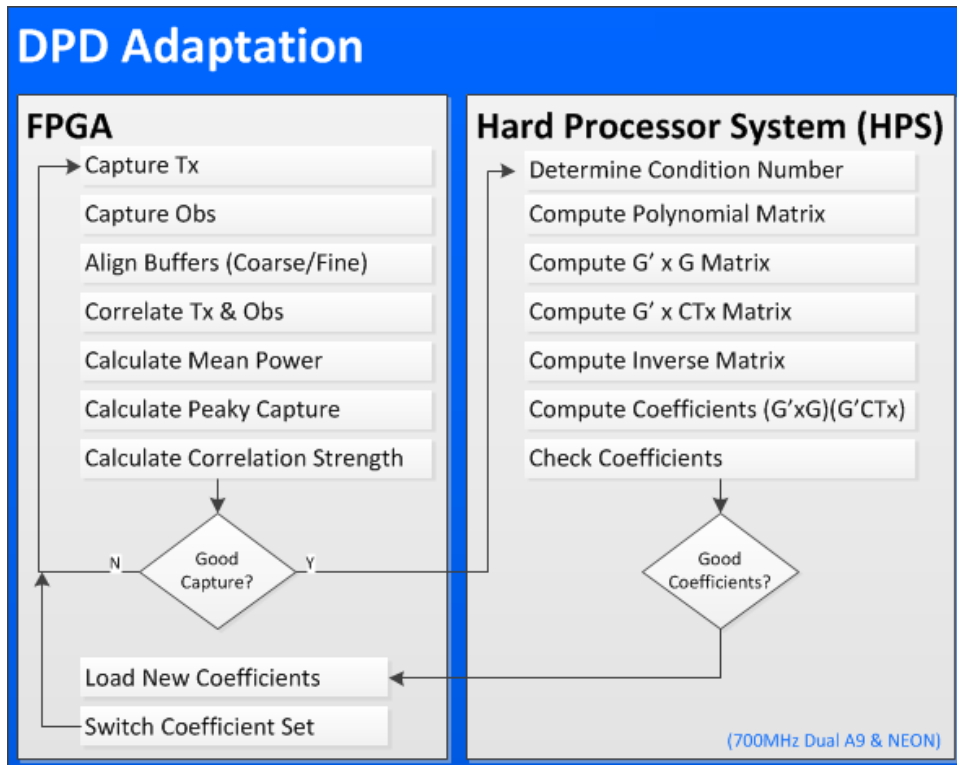
This non-linearity needs to be compensated for by pre-distorting the input signal to be broadcast.

Digital Predistortion is the technology that implements this

This distortion changes with temperature and aging of the PA, so predistortion needs to adapt during operation



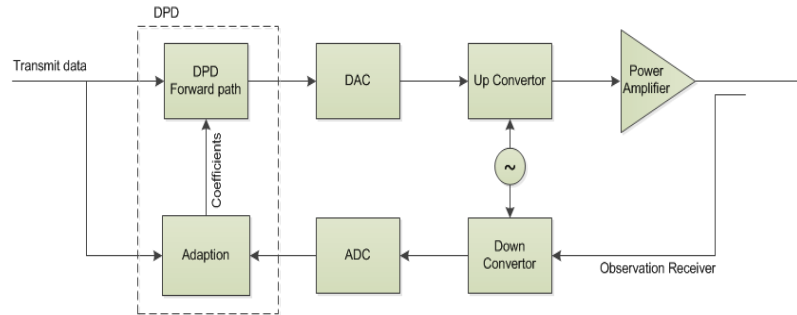
DPD implementation



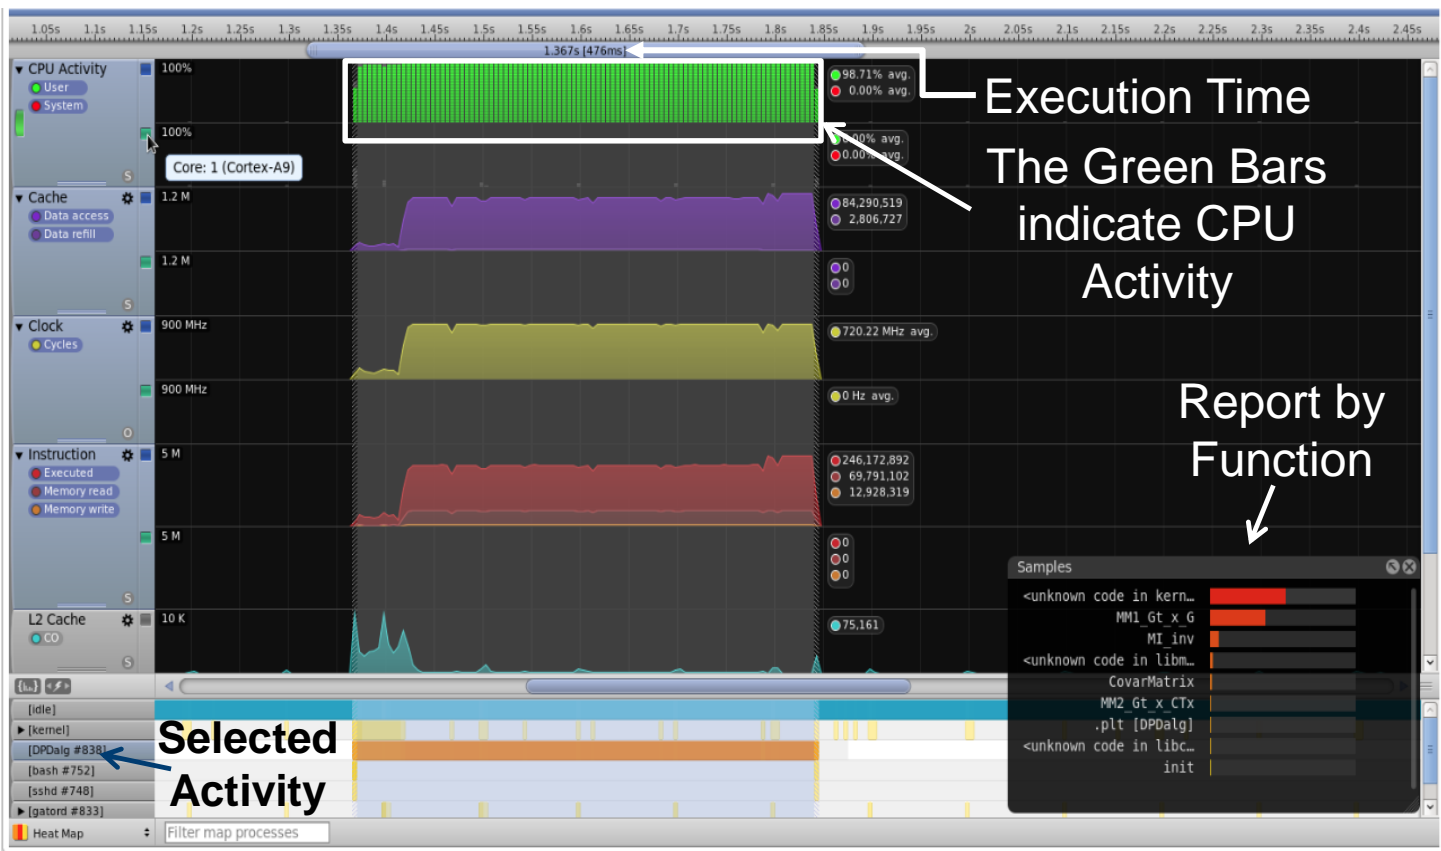
DPD functions

$$\begin{aligned}
 \text{CovarMatrix() } G[36 \times 2040] &= \text{CoRx}[2048 \times 1] \\
 \text{MM1_Gt_x_G()}[36 \times 36] &= G'[36 \times 2040] \times G[2040 \times 36] \\
 \text{MM2_Gt_x_CTx()}[36 \times 1] &= G'[36 \times 2040] \times \text{CTx}[2048 \times 1] \\
 \text{MI_inv()} [36 \times 36] &= \text{INV}(\text{MM1_Gt_x_G}[36 \times 36]) \\
 \text{MI_x_MM2()}[36 \times 1] \text{ Coeffs} &= \text{MM1_Gt_x_G()}[36 \times 36] \times \text{MM2_Gt_x_CTx()}[36 \times 1]
 \end{aligned}$$

$$\text{Coefficients} = \text{INV}(G' \times G) \times (G' \times \text{CTx})$$



Software profiling



Call paths

Self	% Self	Process	% Process	Total	Stack	[Process]/[Thread]/Code	Location
		7,649	100.00%	94.13%	0	[idle]	-
		7,649	100.00%	94.13%	0	{swapper/0 #0}	-
7,649	100.00%	7,649	100.00%	94.13%	0	<unknown code in kernel>	<anonymous>
		467	100.00%	5.75%	0	{DPDalg #838}	-
		467	100.00%	5.75%	0	{DPDalg #838}	-
359	76.87%	359	76.87%	4.42%	160	-MM1_Gt_x_G	DPDalg.c:84
54	11.56%	54	11.56%	0.66%	288	-MI_inv	DPDalg.c:151
19	4.07%	19	4.07%	0.23%	0	<unknown code in libm-2.15.so>	<anonymous>
14	3.00%	14	3.00%	0.17%	416	-CovarMatrix	DPDalg.c:50
9	1.93%	9	1.93%	0.11%	160	-MM2_Gt_x_CTx	DPDalg.c:108
9	1.93%	9	1.93%	0.11%	0	<unknown code in kernel>	<anonymous>
1	0.21%	1	0.21%	0.01%	96	-init	DPDalg.c:27
1	0.21%	1	0.21%	0.01%	0	-.plt [DPDalg]	DPDalg
1	0.21%	1	0.21%	0.01%	0	<unknown code in libc-2.15.so>	<anonymous>
		8	100.00%	0.10%	0	[kernel]	-
		1	100.00%	0.01%	0	[bash #752]	-
		1	100.00%	0.01%	0	[sshd #748]	-
		0	0.00%	0.00%	0	[gator #833]	-
		0	0.00%	0.00%	0	[lighttpd #685]	-

Samples	% Samples	Instances	Function Name	Location
359	76.87%	1	MM1_Gt_x_G	DPDalg.c:84
54	11.56%	1	MI_inv	DPDalg.c:151
19	4.07%	1	<unknown code in libm-2.15.so>	<anonymous>
14	3.00%	1	CovarMatrix	DPDalg.c:50
9	1.93%	1	MM2_Gt_x_CTx	DPDalg.c:108
9	1.93%	1	<unknown code in kernel>	<anonymous>
1	0.21%	1	init	DPDalg.c:27
1	0.21%	1	.plt [DPDalg]	DPDalg
1	0.21%	1	<unknown code in libc-2.15.so>	<anonymous>

Idle
Time

Selected Function

Function wise report

Code view

Samples	% Samples	Line	Source File: /home/vikramb/DPD_WS/DPDalg/DPDalg.c
		83	void MM1 Gt x G()
		84	{
		85	int i,k,j=0,l=0;
		86	float tempi=0, tempr=0;
		87	for (k=0;k<tot;k=(k+length)) //G'xG multiplication logic
		88	{ while(j<tot)
		89	{
29	9.03%	90	for (i=k;i<(k+length);i++)
		91	{
134	41.74%	92	tempr += ((G[i].r)*(G[j].r) - ((G[i].im)*(G[j].im)));
137	42.68%	93	tempi += ((G[i].im)*(G[j].r) + ((G[i].r)*(G[j].im)));
21	6.54%	94	j++;
		95	}
		96	GxG[l].r=tempr;
		97	GxG[l].im=tempi;
		98	tempr=0;tempi=0;l++;
		99	}
		100	j=0;
		101	}

Load by line of 'C' code

Samples	% Samples	Address	Opcode	Disassembly
9	2.80%	0x000087F6	EE777A27	VADD.F32 s15,s14,s15
3	0.93%	0x000087FA	EDC77A00	VSTR s15,[r7,#0]
3	0.93%	0x000087FE	F6455318	MOV r3,#0x5d18
6	1.87%	0x00008802	F2C00302	MOVT r3,#2
3	0.93%	0x00008806	697A	LDR r2,[r7,#0x14]
11	3.43%	0x00008808	0002	LSLS r2,r2,#3
4	1.25%	0x0000880A	189B	ADDS r3,r3,r2
8	2.49%	0x0000880C	3304	ADDS r3,#4
3	0.93%	0x0000880E	ED937A00	VLDR s14,[r3,#0]
3	0.93%	0x00008812	F6455318	MOV r3,#0x5d18
1	0.31%	0x00008816	F2C00302	MOVT r3,#2
1	0.31%	0x0000881A	68FA	LDR r2,[r7,#0xc]
2	0.62%	0x0000881C	0002	LSLS r2,r2,#3
1	0.31%	0x0000881E	189B	ADDS r3,r3,r2
9	2.80%	0x00008820	EDD37A00	VLDR s15,[r3,#0]
13	4.05%	0x00008824	EE277A27	VMUL.F32 s14,s15
6	1.87%	0x00008828	F6455318	MOV r3,#0x5d18
3	0.93%	0x0000882C	F2C00302	MOVT r3,#2
5	1.56%	0x00008830	697A	LDR r2,[r7,#0x14]


Disassembler View

Identify instruction types

Step 1: Latency vs. precision

Comparison of full floating point to single & double-precision

Determine Error (%) and latency trade-offs.



Functions	Single precision in all functions			Double precision in all functions			Function Name	Execution Time (ms)		
	Error (%)		Latency (ms)	Error (%)		Latency (ms)		Double precision	Selected precision	
	I	Q		I	Q					
CovarMatrix()	<10e-05	<10e-05	27	<10e-10	<10e-10	34	CovarMatrix()	34	27	
MM1_Gt_x_G()	<10e-05	<10e-05	321	<10e-10	<10e-10	338	MM1_Gt_x_G()	338	321	
MM2_Gt_x_CTx()	<10e-05	<10e-05	21	<10e-10	<10e-10	29	MM2_Gt_x_CTx()	29	21	
MI_inv()	0.0435	0.019	53	0.00025	0.0003	65	MI_inv()	65	58	
MI_x_MM2()	0%	0%	6	0%	0%	10	MI_x_MM2()	10	6	
Total			428	Total			476	TOTAL	476 ms	433 ms

Step 2: Algorithmic optimization

A matrix multiplied with its transpose results in a square symmetric matrix

This property helps us reduce almost half the number of complex multiplications

All the elements in the lower triangular matrix can be obtained from upper triangular matrix

Function Name	Execution time (ms)	
	Before algorithmic optimization	After algorithmic optimization
CovarMatrix()	27	16
MM1_Gt_x_G()	321	262
MM2_Gt_x_CTx()	21	21
MI_inv()	58	55
MI_x_MM2()	6	6
Total	433 ms	360 ms

$$\begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{12} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \cdots & a_{nn} \end{bmatrix}.$$

Optimized code

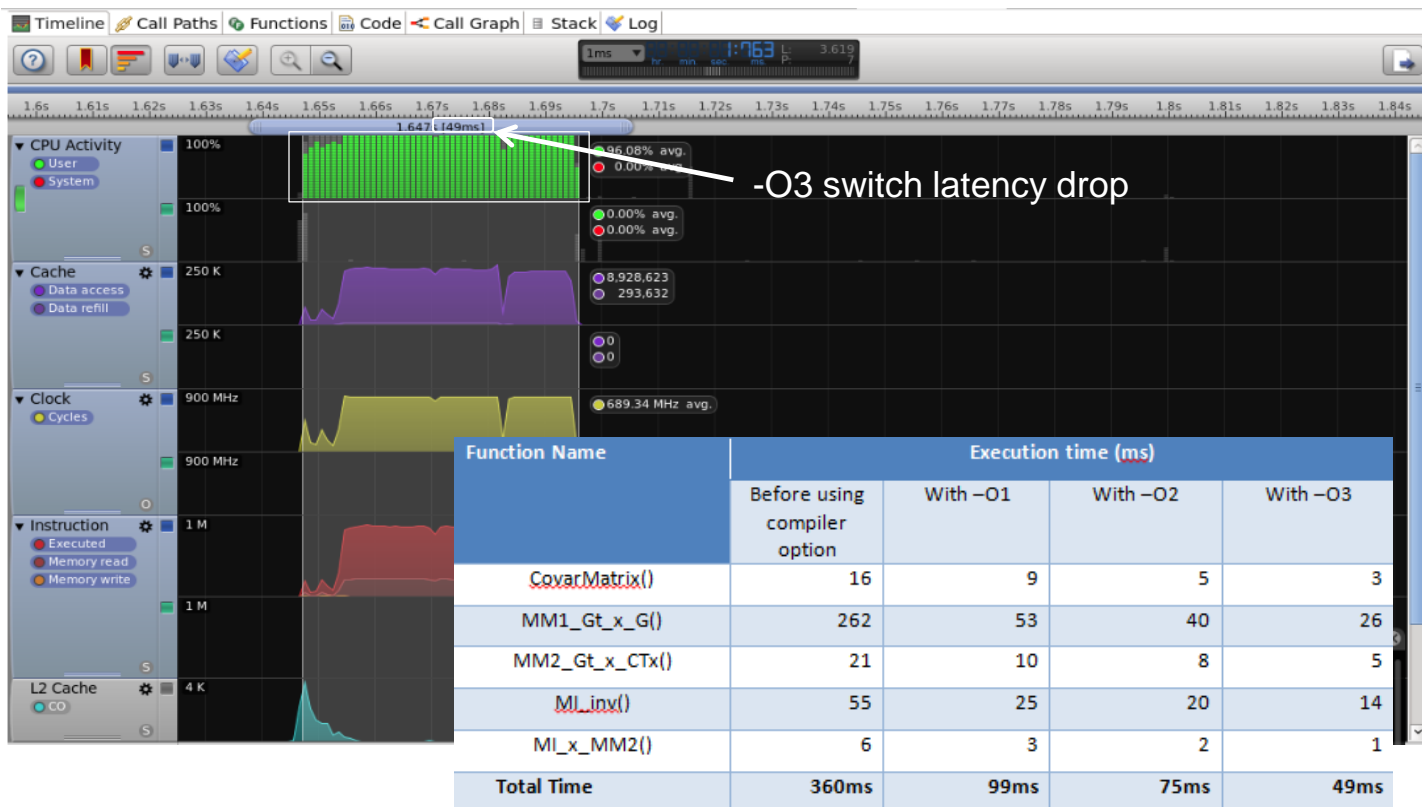
Original Pseudo Code

```
for (k=0;k<tot;k(k+length))
  for(i=k;i<(k+length);i++)
    { tempr +=((G[i].r)*(G[j].r) - (G[i].im)*(G[j].im));
      tempi +=((G[i].im)*(G[j].r) + (G[i].r)*(G[j].r));
```

Optimized Code

```
for (h=0;h<MD;h++)
  for(i=0;i<4;i++)
    for(k=0;l<MD;k++)
      for(l=0;l<4;l++)
        for(m=(MD-k-1);m<(corxlength-k);m++)
          105|37% tempr += ((G+j+(i*corxlength))->r*(G+m+(l*corxlength))->r) -
              ((G+j+(i*corxlength))->im*(G*m+(l*corxlength))->r));
          115|39% tempi += ((G+j+(i*corxlength))->im*(G+m+(l*corxlength))->r) +
              ((G+j+(i*corxlength))->r*(G*m+(l*corxlength))->r));
```

Step 3: Compiler options – OPT-3



Step 4: NEON* vectorization




Data types

- signed/unsigned 8-bit, 16-bit, 32-bit, single-precision

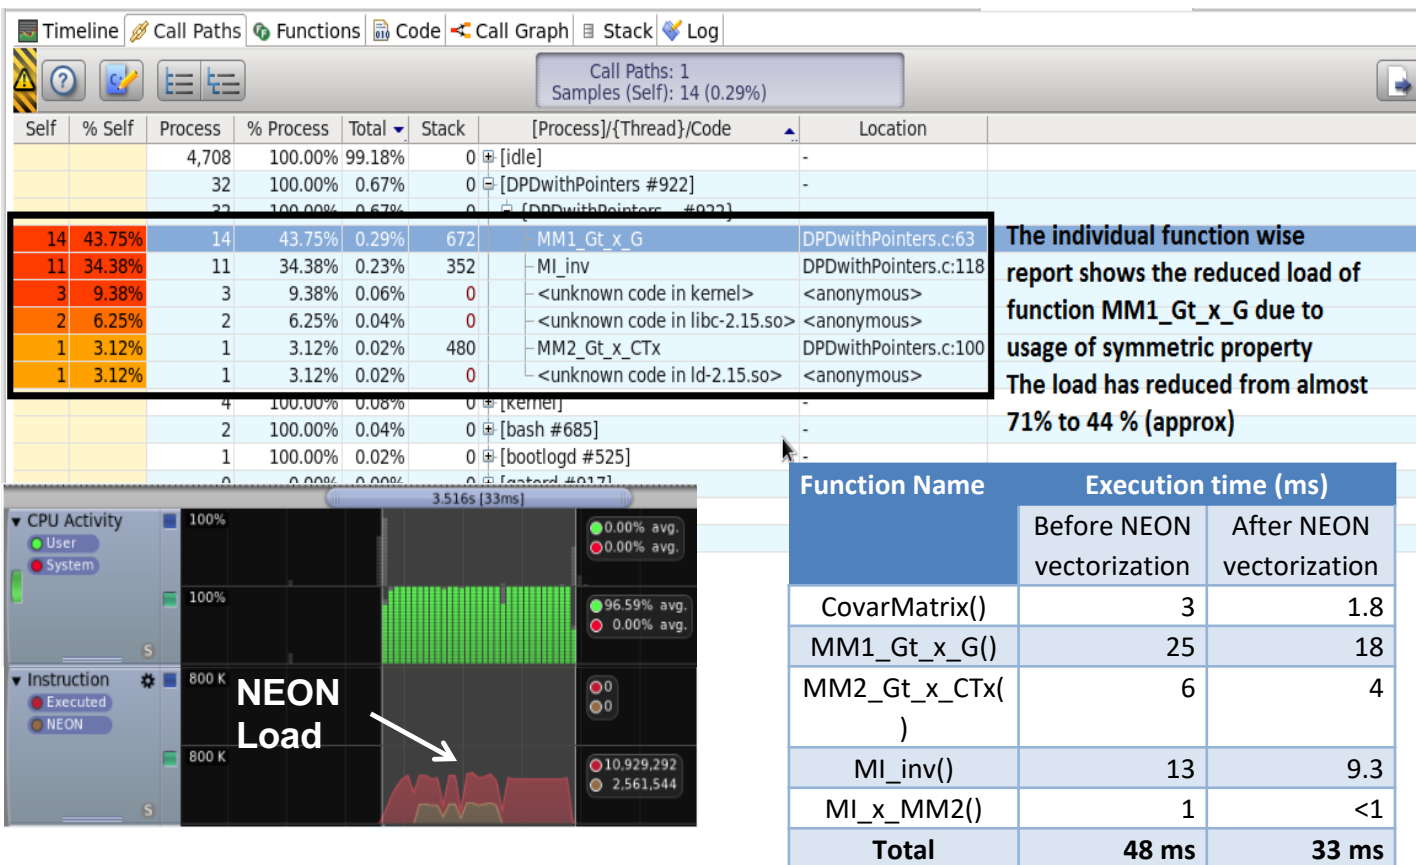
Code Changes

- Use of 1D array instead of 2D array
- Use of keyword 'restrict' for unaligned pointers
- Use of 'const *' for aligned pointers

Description	Resource	Path	Location
▼ i Infos (6 items)			
i vectorized 0 loops in function.	DPDalg.c	/DPDalg	line 26
i vectorized 1 loops in function.	DPDalg.c	/DPDalg	line 83
i vectorized 1 loops in function.	DPDalg.c	/DPDalg	line 107
i vectorized 1 loops in function.	DPDalg.c	/DPDalg	line 128
i vectorized 2 loops in function.	DPDalg.c	/DPDalg	line 49
i vectorized 4 loops in function.	DPDalg.c	/DPDalg	line 150



Vectorization



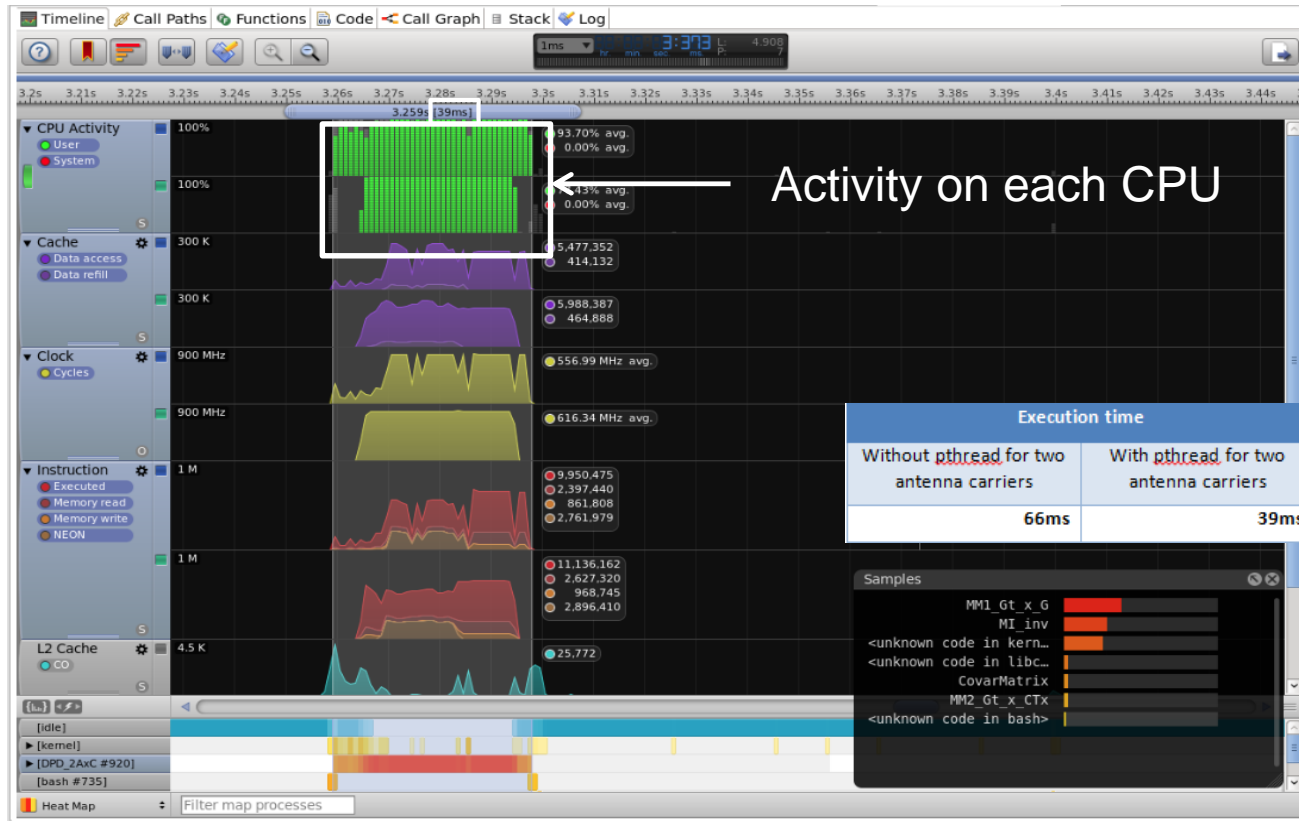
Step 5: Multi-threading

```
#include <pthread.h>
#include <stdio.h>

pthread_t thread1, thread2;
int main()
{
    int rc;
    rc = pthread_create(&thread1, NULL, dpd_algorithm, (void*)&dpd1); //Ax 1
    if (rc) {
        exit(-1);
    }
    rc = pthread_create(&thread2, NULL, dpd_algorithm, (void*)&dpd2); //Ax 2
    if (rc) {
        exit(-1);
    }
    pthread_exit(NULL);
    return 0;
}
```

Self	% Self	Process	% Process	Total	Stack	[Process]/[Thread]/Code	Location
		9,716	100.00%	99.28%	0	[idle]	-
		63	100.00%	0.64%	0	[DPD_2AxC #920]	-
		30	47.62%	0.31%	0	{DPD_2AxC #921}	-
14	22.22%	14	22.22%	0.14%	640	-MM1_Gt_x_G	DPD_2AxC.c:85
11	17.46%	11	17.46%	0.11%	320	-MI_inv	DPD_2AxC.c:156
2	3.17%	2	3.17%	0.02%	0	<unknown code in kernel>	<anonymous>
1	1.59%	1	1.59%	0.01%	128	-CovarMatrix	DPD_2AxC.c:60
1	1.59%	1	1.59%	0.01%	480	-MM2_Gt_x_CT_x	DPD_2AxC.c:135
1	1.59%	1	1.59%	0.01%	0	<unknown code in libc-2.15.so>	<anonymous>
		29	46.03%	0.30%	0	{DPD_2AxC #922}	-
15	23.81%	15	23.81%	0.15%	640	-MM1_Gt_x_G	DPD_2AxC.c:85
11	17.46%	11	17.46%	0.11%	320	-MI_inv	DPD_2AxC.c:156
1	1.59%	1	1.59%	0.01%	128	-CovarMatrix	DPD_2AxC.c:60
1	1.59%	1	1.59%	0.01%	480	-MM2_Gt_x_CT_x	DPD_2AxC.c:135
1	1.59%	1	1.59%	0.01%	0	<unknown code in libc-2.15.so>	<anonymous>

Multi-threading



Summary of SW optimisation techniques

1. Compiler settings
2. Algorithmic optimisation (code architecture modification)
3. Data flow optimisation (code architecture modification)
4. Target processor special hardware (e.g. NEON framework vectorization)
5. Implement Multi-threading/multi-core


Acceleration technique #1:

Bump in the wire / pre-processing / post-processing

Bump in the wire techniques

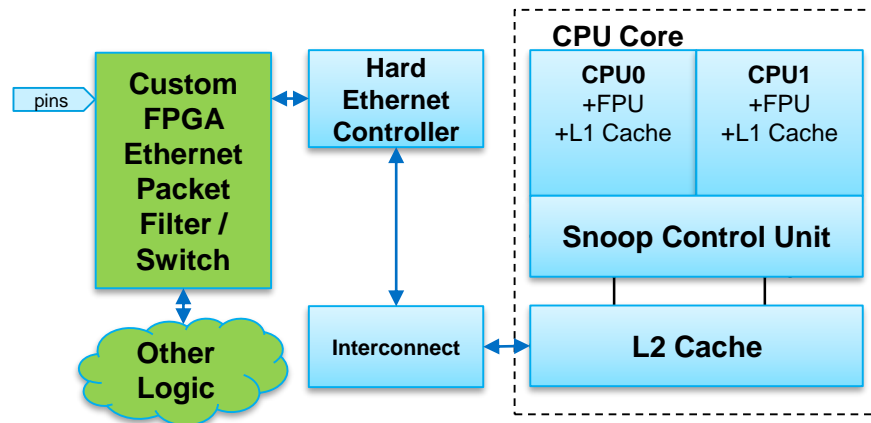
Pre-processing / Post-processing

- Data will be transformed prior / after arrival at the CPU System



Bump in the wire

- The data will be filtered / worked on before arriving at the original interface and can be worked on after data exits the processor



Bump in the wire - potential for deceleration

Pre- or post-processing of data could

Add to the latency between the processor and the I/O.

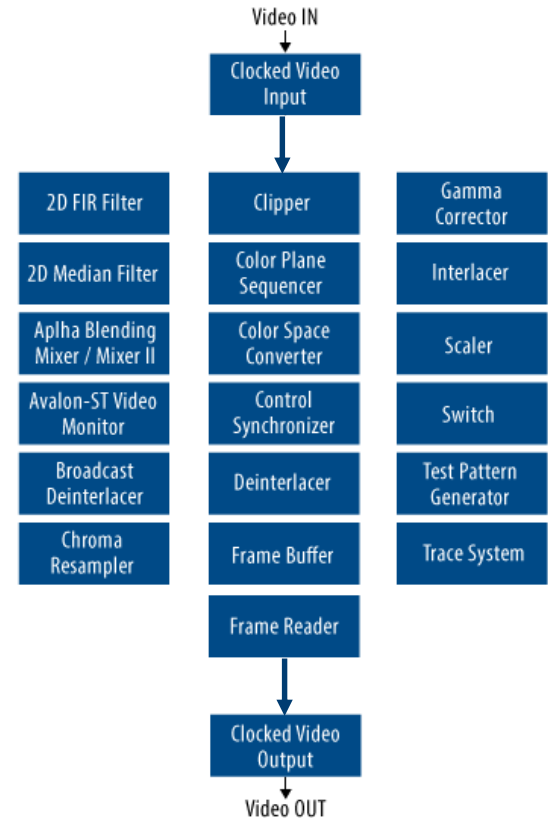
In some instances, the latency of the transaction is critically important to the overall system performance, and needs to be minimized.

Bump in the wire - potential for acceleration

Performance gains should be relatively straight forward to estimate.

CPU offload can significantly speed up memory-intensive functions or reduce CPU utilization

The FPGA is real-time tasks performed with clock-cycle accuracy / determinism.



Acceleration technique #2:

Tightly-coupled instruction set extension



Tightly-coupled instruction set extension

Custom instructions are used to extend the operation of the soft processor

- Data is passed via registers, flags, opcodes to acceleration core from processor

or

- Data is passed via FIFO interface with flags to acceleration core from processor



Note: CI accelerators can have their own local memory

Tightly-coupled instruction set extension - potential for deceleration


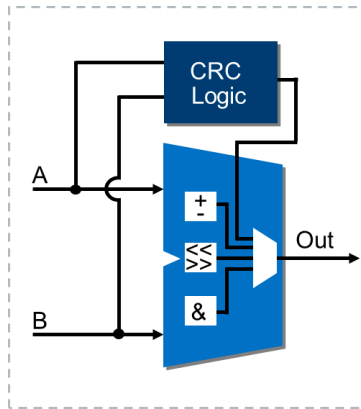
Tightly-coupled instruction set extension could reduce performance by...

- Adding bubbles to the CPU pipeline if the instruction is multi-cycle**
- Slow down the FPGA if timing is hard to close**
- Halt the processor for many cycles if data is fetched externally**

In architectures where the custom instruction can cause a CPU stall, it is important to ensure that efficiency of the CPU pipeline is maintained. If the custom instruction has a dependency on an external data source it could stall for a long time waiting for a new data value.

Tightly-coupled instruction set extension - potential for acceleration

If the accelerated function would be used frequently by critical code in the system, then an acceleration can be achieved. Tight code loops can be accelerated compared to a multi-instruction approach due to the close proximity of the accelerator to the CPU (In or close to the instruction pipeline).



Function	Operation	Cycles to Execute
Arithmetic	add/subtract	5
	multiply	4
	divide	16
	negate/absolute	1
Roots	square root	8
Conversions	integer to float	4
	float to integer	2
Comparisons	min/max	1
	less than/equal	1
	greater than/equal	1
	equal/not equal	1

Example Floating point offload performance

CRC algorithm design example: <https://www.altera.com/support/support-resources/design-examples/intellectual-property/embedded/nios-ii/exm-custom-instruction.html>
Floating point Cycle counts as documented in Intel Platform Designer version 17.1

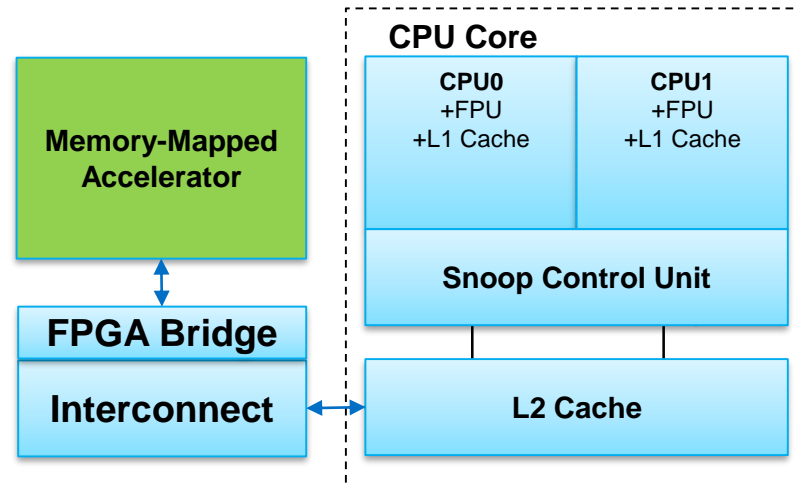
Acceleration technique #3:

Tightly-coupled memory-mapped accelerator

Tightly-coupled memory-mapped accelerator

Accelerator is created on-chip using FPGA resources

- Access to the accelerator is via on-chip bus - Eg. AXI, Avalon
- Parallel, high-bandwidth access is available from processor system to accelerator for read/write access



Tightly-coupled memory-mapped accelerator - potential for deceleration

A tightly-coupled memory-mapped accelerator could reduce performance by...

Being a victim of increased latency

- If the control loop is tight, then the latency can kill performance
- Latency from a SoC hard processor to FPGA can be tens of CPU clock cycles


How do we actually measure latency?

Example scenario:

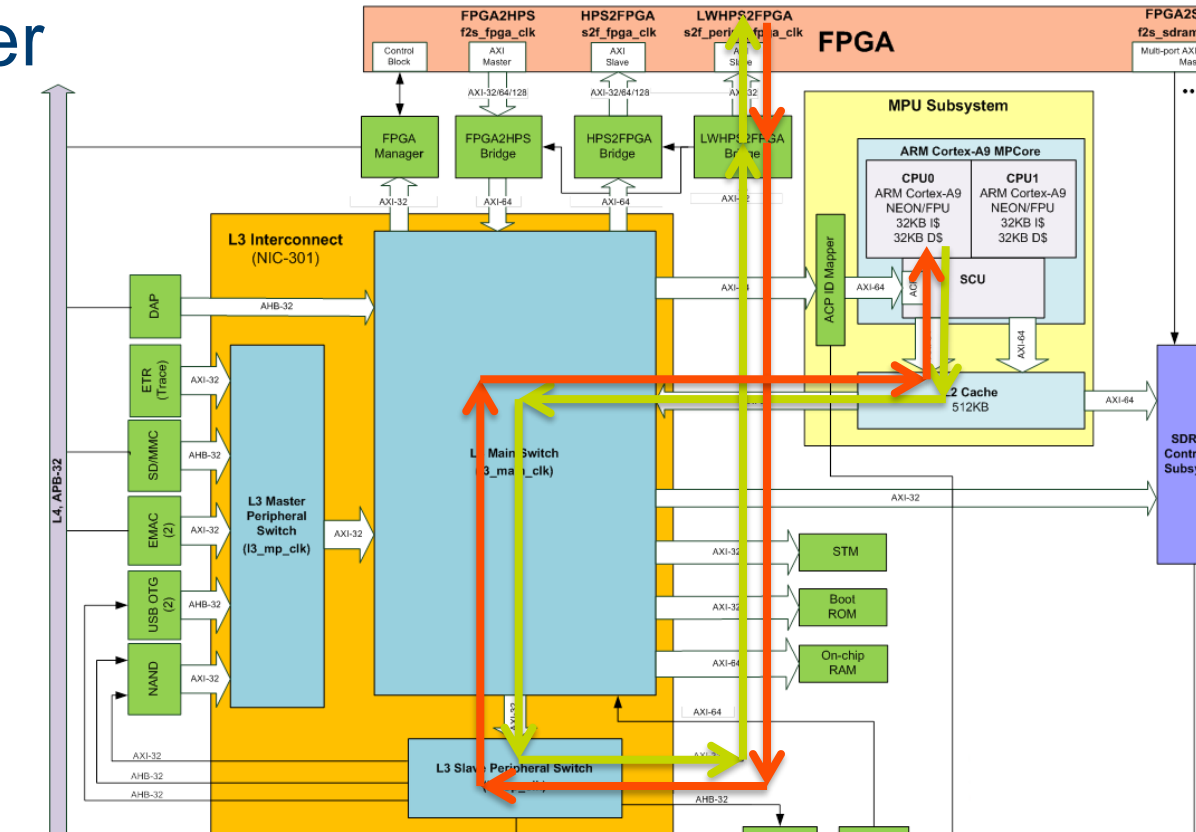
- What is the latency of a 12-byte read?
- What is the latency of a 48-byte read?
- What is the latency of a 20-byte write?

[illegible]

Clarification - what are you trying to do?




So really, we can benchmark a read and write together



Example latency measurement technique

- Software debugger can timestamp cycle counts for a series of transactions
- Hardware or FPGA logic analyzer can measure period between loops of transactions

Cycles		Detail
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
568	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
568	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
560	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
568	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
560	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4
568	LDM	r0, {r1-r3}
0	STM	r0, {r1-r5}
0	B	asm_test3 ; 0x10002F4



Tightly-coupled memory-mapped accelerator - potential for acceleration

- Streaming data fits well with memory-mapped accelerators
- Best performance if results can be read separately to data coming in - the CPU can work on other processing in the mean-time
- DMA can further offload processor from reading directly from peripheral

DPD example - FPGA acceleration

MM1_Gt_x_G() calculation consumes 50% of the entire processing time

- 18ms/39ms (2 antennas)

FPGA acceleration can reduce this time to **<<0.1ms**

- Also exploit matrix shifts (rows are delayed versions)

Total Latency 22ms

- for 2 antennas on low-cost FPGA

DPD Example - Optimization Summary

		Cores used (Cortex A9)	Optimization	Time (ms) per Antenna
Software Optimization	{	1	OPT-0 (Vanilla implementation)	433
		1	OPT-0 (algorithmic optimization)	360
		1	OPT-1	99
		1	OPT-2	75
		1	OPT-3	48
		1	OPT-3 + NEON Vectorization	33
		2	OPT-3 + NEON + pthread (for two antennas)	39/2
Hardware acceleration	{	2	OPT-3 + NEON + pthread (2 antennas) + FPGA Acceleration	21/2

~2x FPGA HW Acceleration vs Optimised Software

All numbers were measured on a DPD implementation targeting Intel Cyclone V SoC. More detail and full figures available at https://www.altera.com/en_US/pdfs/literature/wp/wp-01248-dpd-profiling-and-optimization-with-altera-socs.pdf


Acceleration technique #4:

External bus-connected accelerator card

External, bus-connected Accelerator card

CPU is connected to accelerator via external link, E.g. PCI Express*

- Accelerator is memory-mapped to CPU
- Interrupts also available across the link
- Multiple accelerators may be installed in the same system



Accelerator card - potential for deceleration


An accelerator card could reduce performance by...

Being a victim of increased latency

- If the control loop is tight, then the latency can kill performance
- Latency from a CPU to a PCIe* card can be much longer than on-chip interfaces

External bus-connected accelerator card - potential for acceleration

Where data can be constantly streamed over the interface to the accelerator card, and resultant calculations returned, the latency of the link only serves as an initial latency, and once the pipeline is filled, full-bandwidth use can be made of the accelerator.



Database Acceleration Trial -

10x faster Real-time data analytics¹

2x Traditional data warehousing²

3x Storage Compression³

Genomics sequencing Trial -

50x Pair HMM speedup⁴

1. Based on database queries run with SWARM64 acceleration vs. no acceleration. Testing performed by Swarm64. See System Configurations page for more details.
2. Data warehousing tested with queries and data taken from TPC-DS benchmark. Testing performed by Swarm64. See System Configurations for more details.
3. Based on database size run with SWARM64 acceleration vs. no acceleration. Testing performed by Swarm64. See System Configurations page for more details.
4. Based on tests by the Broad Institute. See System Configurations page for more details of Test configuration.

Tests measure performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance.

Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit www.intel.com/benchmarks.


Acceleration technique #5:

Cache-coherent accelerator with on-chip processor interface

Cache-coherent accelerator with on-chip processor interface

Accelerator is created from FPGA on-chip resources

- Access for data transit is via ACP interface to cache-coherent port
- Transactions marked as cacheable will be directly sourced from cache as appropriate



Cache-coherent accelerator - potential for deceleration

An cache-coherent accelerator reduce performance by...

Thrashing the caches

- Large data transfers through cache interfaces will cause existing cache data to be swapped out
- Other CPU workloads could slow down if this data needs to be re-fetched
- Cache line locking can help, but has knock-on effects on performance

Cache-coherent accelerator - potential for acceleration

Participation in cache-coherence can have major performance benefits in some instances compared to memory mapped systems.

- No need for CPU to flush data to external memory
- Direct hardware access to cached data without DDR interfacing
- Payload returned close to CPU, so long as it fits in the cache

Virtualization:

The effects of virtualization on hardware performance

Use of virtualization is increasingly popular

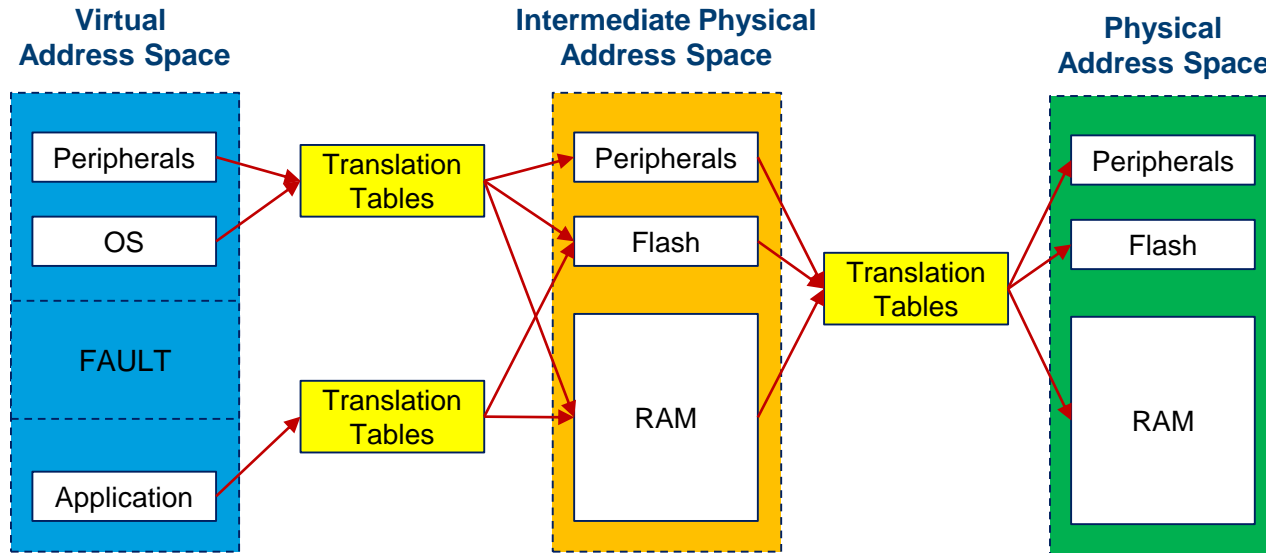
Multi-user environments

More than one O/S required to run legacy and new code

Hypervisors used to manage resources and permissions to hardware

Additional architectural complexity to support hypervisor architecture

Virtualization - adds additional address decode for each virtual OS.



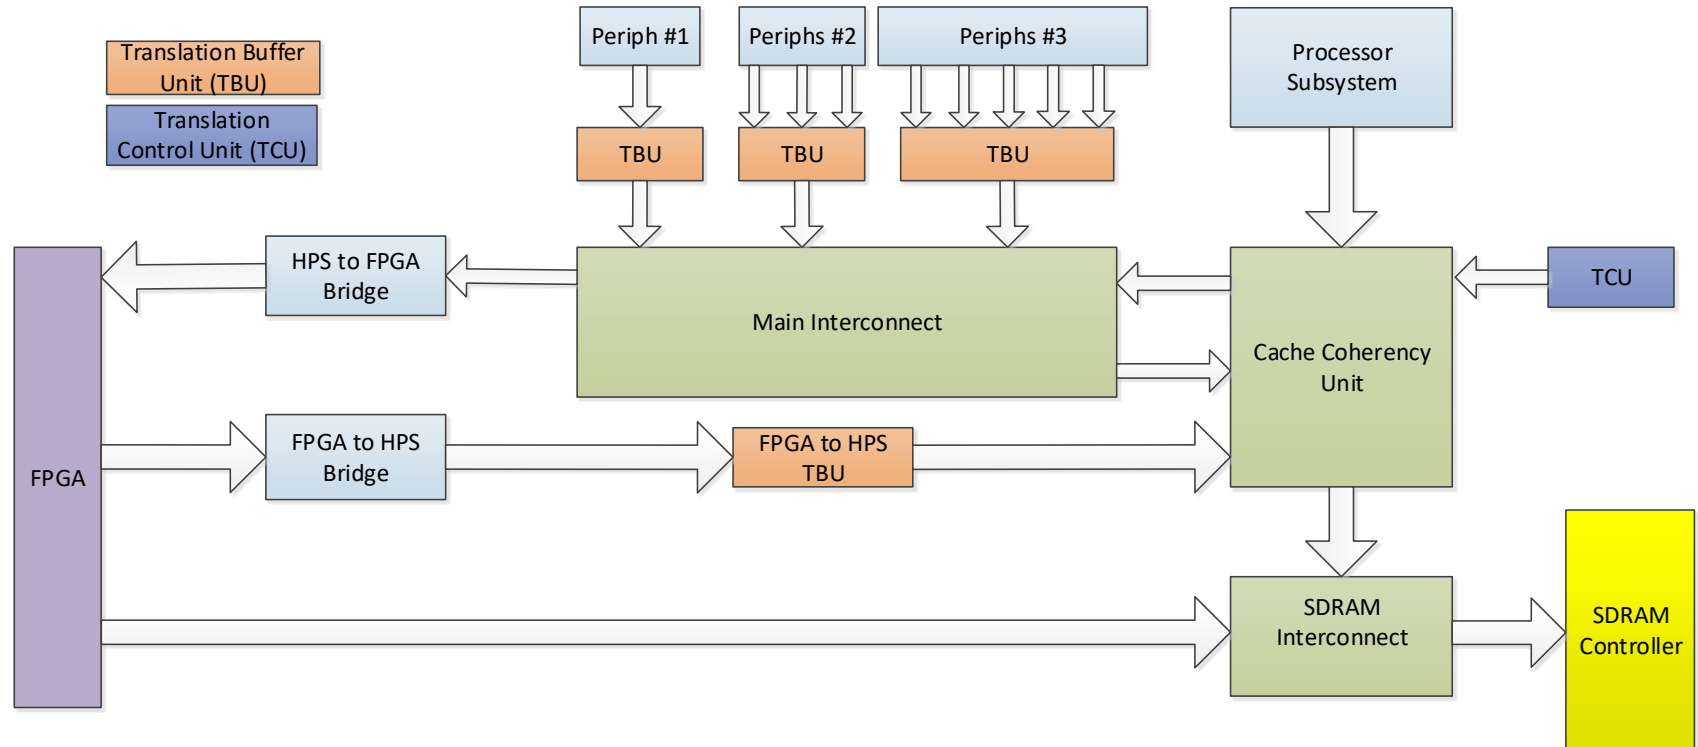
Virtualization - potential for deceleration

Virtualization could decrease system performance by ...

Adding additional address decode logic

- Architecture is key here. Local caching of translation of translation buffers for the System MMU is key to maintain performance.

Example System MMU infrastructure



Virtualization - potential for acceleration

Virtualization could increase system performance by ...

Allowing multiple O/S to be hosted on one device

- Off-chip latency has been seen to be critical in some designs. Memory can be safely shared, resources saved and systems optimized for data throughput.

Summary

Summary

1. Bump in the wire:

Pros: Independent of processor & memory, deterministic, can be low latency

Cons: Adds latency to arrival of data

2. Processor custom instructions

Pros: local data, tight integration with software

Cons: Can stall processor/block other processes

3. On-chip memory-mapped accelerator

Pros: Can act independently of processor, processor can do other tasks, data can stream direct from peripheral

Cons: Adds latency to arrival of data to processor, has to be on-chip

4. External bus-connected accelerator card

Pros: Can act independently of processor, processor can do other tasks, data can stream direct from external interface, easy to implement, can support multiple instances

Cons: Latency is high

5. Cache-coherent accelerator

Pros: Local data, low latency, potential high performance

Cons: Data set must be small and cache aligned, processor/data has to be managed to avoid cache thrash

Conclusion

The performance improvement of a processor system with accelerators is greatly influenced by **architectural choices**.

The architect needs to consider **latency** and **data flow** dependencies to make sound implementation choices. In that way pipeline bubbles and other effects that can negatively affect performance are minimized.

Swarm64 system configuration: Supermicro* SuperServer 2028U-TR4+ with Super X10DRU-i+ Mainboard, 2X Intel® Xeon® E5-2695 v4 CPUs, 8X Samsung* 32GB DDR4-2400 ECC RAM.

Note: This is SQL to relational database, not SQL to semi/unstructured data. Data Warehousing tested with data taken from the TPC-DS benchmarking suite and tested against the 99 TPC-DS query templates.

Broad Institute system configuration: Intel® Xeon® processor E5-2697 v2 @ 2.70 GHz, 2 sockets/12 cores per socket, 128 GB RAM, 2 TB Seagate HDD ST2000DM001 with Intel Arria 10 GX FPGA compared to Intel® Xeon® processor E5-2699 v4 @ 2.20 GHz, 2 sockets/22 cores per socket, 256 GB RAM, 2 TB Intel® SSD Data Center P3700 Series with single-core Intel® Advanced Vector Extensions (AVX) technology.

Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit www.intel.com/benchmarks

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. No computer system can be absolutely secure. Check with your system manufacturer or retailer or learn more at intel.com.

Optimization Notice: Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice Revision #20110804

Intel, the Intel logo, and neon are trademarks of Intel Corporation or its subsidiaries in the U.S. and/or other countries.

*Other names and brands may be claimed as the property of others.

© Intel Corporation.

